

STAI 30 - Class #2:

Nathan Taback

2018-01-15

Today's Class

- Histograms and density functions
- Statistical data
- Data wrangling
- Tidy data

Histograms and Density Functions

Histograms and Density Functions

- The histogram of a variable is a graphical method to visualize the distribution of a single variable.
- To construct a basic histogram:
 1. Divide the data into intervals (called bins).
 2. Count the number of observations that are contained in the bin.
 3. Plot rectangles with height equal to the count from (2) and width equal to the width of the bin.

Histograms and Density Functions


- Different bin width will yield different histograms

```
p1 <- ggplot(data = happinessdata2016, aes(x = life_ladder)) +  
  geom_histogram(binwidth = 0.1)  
p2 <- ggplot(data = happinessdata2016, aes(x = life_ladder)) +  
  geom_histogram(binwidth = 1.0)  
  
ggplot_build(p2)
```



```
## $data  
## $data[[1]]  
##      y count  x xmin xmax  density  ncount  ndensity PANEL group ymin  
## 1  8      8  3  2.5  3.5 0.05673759 0.1904762 26.85714     1    -1     0  
## 2 25     25  4  3.5  4.5 0.17730496 0.5952381 83.92857     1    -1     0  
## 3 42     42  5  4.5  5.5 0.29787234 1.0000000 141.00000     1    -1     0  
## 4 41     41  6  5.5  6.5 0.29078014 0.9761905 137.64286     1    -1     0  
## 5 20     20  7  6.5  7.5 0.14184397 0.4761905 67.14286     1    -1     0  
## 6  5      5  8  7.5  8.5 0.03546099 0.1190476 16.78571     1    -1     0  
##      ymax colour  fill size linetype alpha  
## 1      8      NA grey35 0.5      1      NA  
## 2     25      NA grey35 0.5      1      NA  
## 3     42      NA grey35 0.5      1      NA  
## 4     41      NA grey35 0.5      1      NA  
## 5     20      NA grey35 0.5      1      NA  
## 6      5      NA grey35 0.5      1      NA  
##  
##  
## $layout  
## <ggproto object: Class Layout>
```

```
## facet: <ggproto object: Class FacetNull, Facet>
##   compute_layout: function
##   draw_back: function
##   draw_front: function
##   draw_labels: function
##   draw_panels: function
##   finish_data: function
##   init_scales: function
##   map: function
##   map_data: function
##   params: list
##   render_back: function
##   render_front: function
##   render_panels: function
##   setup_data: function
##   setup_params: function
##   shrink: TRUE
##   train: function
##   train_positions: function
##   train_scales: function
##   vars: function
##   super: <ggproto object: Class FacetNull, Facet>
## finish_data: function
## get_scales: function
## map: function
## map_position: function
## panel_layout: data.frame
## panel_ranges: list
## panel_scales: list
## render: function
## render_labels: function
## reset_scales: function
## setup: function
## train_position: function
## train_ranges: function
## xlabel: function
```

```
##      ylabel: function
##      super:  <ggproto object: Class Layout>
##
## $plot
```



```
grid.arrange(p1,p2,nrow = 1)
```



Mathematical Definition of Histogram

- The bins of the histogram are the intervals:

$$[x_0 + mh, x_0 + (m + 1)h).$$

x_0 is the origin, $m = \dots, -1, 0, 1, \dots$ indexes the bins, and $h = (x_0 + (m + 1)h) - (x_0 + mh)$ is the bin width.

Example - Mathematical Definition of Histogram

```
dat <- data_frame(x = c(1,2,2.5,3,7))  
dat$x
```

```
[1] 1.0 2.0 2.5 3.0 7.0
```

Let $x_0 = 0.5, h = 0.25, m = 1, \dots, 29$

```
seq(0.5,7.5,by = 0.25)
```

```
[1] 0.50 0.75 1.00 1.25 1.50 1.75 2.00 2.25 2.50 2.75 3.00 3.25 3.50 3.75  
[15] 4.00 4.25 4.50 4.75 5.00 5.25 5.50 5.75 6.00 6.25 6.50 6.75 7.00 7.25  
[29] 7.50
```

The bins are: $[0.50, 0.75), [0.75, 1.00), [1.00, 1.25), \dots, [7.25, 7.50)$.




Example - Mathematical Definition of Histogram

- The bins can be used to construct rectangles with width $h = 0.25$ and height y .
- y will be called density.
- The area of these rectangles is hy .
- We would like the area of these rectangles, hy , to be the same as the proportion of data in the bin. This will make the sum of all areas equal 1.
- Let n be the number of observations. Then,

$$hy = \frac{\#\{X_i \text{ in bin}\}}{n}$$

- In this example, $n = 5$, and $X_1 = 1, X_2 = 2, X_3 = 2.5, X_4 = 3, X_5 = 7$.

Example - Mathematical Definition of Histogram



Mathematical Definition of Histogram

$$\hat{f}(x) = \frac{1}{hn} \#\{X_i \text{ in same bin as } x\}$$

is called the **histogram estimator**.

$\hat{f}(x)$ is an estimate of the density at a point x .


To construct the histogram we have to choose an origin x_0 and bin width h .

Choosing Origin and Bin Width in R


Same bin width but different origin

```
dat <- data_frame(x = c(1,2,2.5,3,7))
p0.5 <- ggplot(data = dat, aes(x = x)) +
  geom_histogram(breaks = seq(0.5,7.5,.25)) +
  labs(subtitle = "origin = 0.5, bin width = 0.25")
p2 <- ggplot(data = dat, aes(x = x)) +
  geom_histogram(breaks = seq(2,7.5,.25)) +
  labs(subtitle = "origin = 2.0, bin width = 0.25")
grid.arrange(p0.5,p2,nrow = 1)
```

origin = 0.5, bin width = 0.25



origin = 2.0, bin width = 0.25



Statistical data

What is statistical data?

- Statistical data is obtained by observing (random) variables.
- A random variable can be given a precise mathematical definition that we will cover later in the course.
- In this class we will discuss examples.

Observing a few variables on STAI30 students

- What is your height?
- How many years have been at UofT?
- What is your sex (male or female)?

Collecting this data will generate three variables: height, years, and sex.

Enter variables on STA130 students

```
height <- c()  
years <- c()  
sex <- c()
```

Put the variables into an R data frame.

NB: `data_frame` is the tidyverse version of base R `data.frame`.

```
sta130_dat <- data_frame(height, years, sex)
```

We could have entered this in a spreadsheet program like MS Excel, saved it as a CSV file, then imported the file into R.

Tidy data

Tidy data

There are three interrelated rules which make a dataset tidy:

1. Each variable must have its own column.
2. Each observation must have its own row.
3. Each value must have its own cell.

```
## # A tibble: 6 x 4
##   country    year cases population
##   <chr>    <int> <int>      <int>
## 1 Afghanistan 1999     745   19987071
## 2 Afghanistan 2000    2666   20595360
## 3      Brazil 1999   37737   172006362
## 4      Brazil 2000   80488   174504898
## 5        China 1999  212258  1272915272
## 6        China 2000  213766  1280428583
```

```
## # A tibble: 6 x 3
##   country    year      rate
##   *      <chr> <int>    <chr>
## 1 Afghanistan 1999    745/19987071
## 2 Afghanistan 2000   2666/20595360
## 3      Brazil 1999  37737/172006362
## 4      Brazil 2000  80488/174504898
## 5        China 1999 212258/1272915272
## 6        China 2000 213766/1280428583
```


Tidy data

“For a given dataset, it is usually easy to figure out what are observations and what are variables, but it is surprisingly difficult to precisely define variables and observations in general.” (Wickham, 2014)

A general rule of thumb:

- It is easier to describe functional relationships between variables (e.g., z is a linear combination of x and y , density is the ratio of weight to volume) than between rows.
- It is easier to make comparisons between groups of observations (e.g., average of group a vs. average of group b) than between groups of columns.

(Wickham, 2014)

Data wrangling

Data wrangling

- The `ggplot` library implements a **grammar of graphics**.
- Similarly the `dplyr` library presents a grammar for data wrangling.

The Economic Guide to Picking a Major

FiveThirtyEight

Politics Sports Science & Health **Economics** Culture


SEP. 12, 2014 AT 7:37 AM

The Economic Guide To Picking A College Major

By Ben Casselman

Filed under Higher Education

Get the data on GitHub



Students walk across the campus of UCLA in Los Angeles. KEVORK DJANSEZIAN / GETTY IMAGES

“...A college degree is no guarantee of economic success. But through their choice of major, they can take at least some steps toward boosting their odds.”

The Economic Guide to Picking a Major

- The data used in the article is from the American Community Survey 2010-2012 Public Use Microdata Series.
- We can use the `fivethirtyeight` library in R.

Data behind the article

```
# uncomment this next line if fivethirtyeight lib is not installed
# install.packages("fivethirtyeight")
library(fivethirtyeight) # load the library
glimpse(college_recent_grads)
```

```
## Observations: 173
## Variables: 21
## $ rank                <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,...
## $ major_code          <int> 2419, 2416, 2415, 2417, 2405, 2418...
## $ major               <chr> "Petroleum Engineering", "Mining A...
## $ major_category      <chr> "Engineering", "Engineering", "Eng...
## $ total               <int> 2339, 756, 856, 1258, 32260, 2573,...
## $ sample_size         <int> 36, 7, 3, 16, 289, 17, 51, 10, 102...
## $ men                 <int> 2057, 679, 725, 1123, 21239, 2200,...
## $ women               <int> 282, 77, 131, 135, 11021, 373, 960...
## $ sharewomen          <dbl> 0.1205643, 0.1018519, 0.1530374, 0...
## $ employed            <int> 1976, 640, 648, 758, 25694, 1857, ...
## $ employed_fulltime   <int> 1849, 556, 558, 1069, 23170, 2038,...
## $ employed_parttime   <int> 270, 170, 133, 150, 5180, 264, 296...
## $ employed_fulltime_yearround <int> 1207, 388, 340, 692, 16697, 1449, ...
## $ unemployed          <int> 37, 85, 16, 40, 1672, 400, 308, 33...
## $ unemployment_rate   <dbl> 0.018380527, 0.117241379, 0.024096...
## $ p25th               <dbl> 95000, 55000, 50000, 43000, 50000,...
## $ median              <dbl> 110000, 75000, 73000, 70000, 65000...
## $ p75th               <dbl> 125000, 90000, 105000, 80000, 7500...
## $ college_jobs        <int> 1534, 350, 456, 529, 18314, 1142, ...
## $ non_college_jobs    <int> 364, 257, 176, 102, 4440, 657, 314...
## $ low_wage_jobs       <int> 193, 50, 0, 0, 972, 244, 259, 220,...
```


Select variables/columns using select ()

To retrieve a data frame with only major, number of male and female graduates we use the select () function in the dplyr library.

```
select(college_recent_grads, major, men, women)
```

```
## # A tibble: 173 x 3
##               major    men women
##               <chr> <int> <int>
## 1      Petroleum Engineering    2057    282
## 2 Mining And Mineral Engineering     679     77
## 3      Metallurgical Engineering     725    131
## 4 Naval Architecture And Marine Engineering    1123    135
## 5      Chemical Engineering   21239   11021
## 6      Nuclear Engineering     2200     373
## 7      Actuarial Science        832     960
## 8 Astronomy And Astrophysics     2110    1667
## 9      Mechanical Engineering   12953    2105
## 10 Electrical Engineering     8407    6548
## # ... with 163 more rows
```

Select observations/rows using `filter()`

If we want to retrieve only those observations (rows) that pertain to engineering majors then we need to specify that the value of the major variable is Electrical Engineering.

```
EE <- filter(college_recent_grads, major == "Electrical Engineering")
glimpse(EE)
```

```
## Observations: 1
## Variables: 21
## $ rank          <int> 10
## $ major_code     <int> 2408
## $ major          <chr> "Electrical Engineering"
## $ major_category <chr> "Engineering"
## $ total          <int> 81527
## $ sample_size    <int> 631
## $ men            <int> 8407
## $ women          <int> 6548
## $ sharewomen     <dbl> 0.4378469
## $ employed       <int> 61928
## $ employed_fulltime <int> 55450
## $ employed_parttime <int> 12695
## $ employed_fulltime_yearround <int> 41413
## $ unemployed     <int> 3895
## $ unemployment_rate <dbl> 0.05917385
## $ p25th          <dbl> 45000
## $ median         <dbl> 60000
## $ p75th          <dbl> 72000
```

```
## $ college_jobs      <int> 45829  
## $ non_college_jobs  <int> 10874  
## $ low_wage_jobs     <int> 3170
```

NB: == is a test for equality and is different than =.

Combine `select()` and `filter()`

- We can drill down to get certain pieces of information using `filter()` and `select()` together.
- The median variable is median salary.

```
select(filter(college_recent_grads, median >= 60000), major, men, women)
```

(1) Which students, and (2) variables are in this data frame?



Respond at **PollEv.com/nathantaback**



Text **NATHANTABACK** to **37607** once to join, then **A, B, C, D, or E**

(1) All students in the original data set; (2) all variables in the data set.

A

(1) All students in the original data set in a major where the median

D

salary is at most than 60,000; (2) all variables in the data set.

D

(1) All students in the original data set in a major where the median salary is at least than 60,000; (2) all variables in the data set.

C

(1) All students in the original data set in a major where the median salary is at least than 60,000; (2) three variables: major, men, women

D

(1) All students in the original data set in a major where the median salary is at least than 60,000; (2) all variables in the data set.

E

The pipe operator %>%

In the code:

```
select(filter(college_recent_grads, median >= 60000), major, men, women)
```

filter is nested inside select.

The pipe operator allows is an alternative to nesting and yields easier to read code. The same expression can be written with the pipe operator

```
college_recent_grads %>%  
  filter(median >= 60000) %>%  
  select(major, men, women)
```

Create new variables from existing variables using mutate()

What percentage of graduates from each major where the median earnings is at least \$60,000 are men ?

```
college_recent_grads %>%  
  filter(median >= 60000) %>%  
  select(major, men, women) %>%  
  mutate(total = men + women, pct_male = round((men / total)*100, 2))
```

Compare to nested code:

```
mutate(select(filter(college_recent_grads, median >= 60000),  
            major, men, women),  
       total = men + women, pct_male = round((men / total)*100, 2))
```

Create new variables from existing variables using mutate()

major	men	women	total	pct_male
Petroleum Engineering	2057	282	2339	87.94
Mining And Mineral Engineering	679	77	756	89.81
Metallurgical Engineering	725	131	856	84.70
Naval Architecture And Marine Engineering	1123	135	1258	89.27
Chemical Engineering	21239	11021	32260	65.84
Nuclear Engineering	2200	373	2573	85.50
Actuarial Science	832	960	1792	46.43
Astronomy And Astrophysics	2110	1667	3777	55.86
Mechanical Engineering	12953	2105	15058	86.02
Electrical Engineering	8407	6548	14955	56.22
Computer Engineering	33258	8284	41542	80.06

Aerospace Engineering	65511	16016	81527	80.35
Biomedical Engineering	80320	10907	91227	88.04
Materials Science	2949	1330	4279	68.92

Create new variables from existing variables using mutate()

- Suppose that we would like to create a categorical variable to identify majors with 45% and 55% women (ie., approximately equal numbers of males and females).
- We can use ifelse() in a mutate() statement.

```
college_recent_grads %>%  
  select(major, men, women) %>%  
  mutate(total = men + women, pct_female = round((women / total)*100, 2),  
         male.bias = ifelse(pct_female >= 45 & pct_female <= 55, "No", "Yes")) %>%  
  select(major, male.bias)
```

```
## # A tibble: 173 x 2  
##           major male.bias  
##           <chr>      <chr>  
## 1      Petroleum Engineering      Yes  
## 2 Mining And Mineral Engineering      Yes  
## 3 Metallurgical Engineering      Yes  
## 4 Naval Architecture And Marine Engineering      Yes  
## 5      Chemical Engineering      Yes  
## 6      Nuclear Engineering      Yes  
## 7      Actuarial Science          No  
## 8 Astronomy And Astrophysics      Yes  
## 9      Mechanical Engineering      Yes  
## 10 Electrical Engineering      Yes
```

... with 163 more rows

Rename variables using rename ()

- It's considered bad practice in R to use periods in variable names.
- We can use rename () to change the name of sex.equal to sex_equal.

```
my_college_dat <- college_recent_grads %>%  
  select(major, men, women, median) %>%  
  mutate(total = men + women, pct_female = round((women / total)*100, 2),  
         sex.equal = ifelse(pct_female >= 45 & pct_female <= 55, "No", "Yes")) %>%  
  select(major, sex.equal, median)  
  
my_college_dat <- my_college_dat %>%  
  rename(sex_equal = sex.equal, salary_median = median)  
glimpse(my_college_dat)
```

```
## Observations: 173  
## Variables: 3  
## $ major      <chr> "Petroleum Engineering", "Mining And Mineral Eng...  
## $ sex_equal   <chr> "Yes", "Yes", "Yes", "Yes", "Yes", "Yes", "No", ...  
## $ salary_median <dbl> 110000, 75000, 73000, 70000, 65000, 65000, 62000...
```

Sort a data frame using arrange ()

```
my_college_dat %>%  
  arrange(salary_median) %>%  
  select(major, salary_median) %>%  
  arrange(desc(salary_median))
```

```
## # A tibble: 173 x 2  
##  
##           major salary_median  
##           <chr>         <dbl>  
## 1      Petroleum Engineering    110000  
## 2 Mining And Mineral Engineering    75000  
## 3 Metallurgical Engineering    73000  
## 4 Naval Architecture And Marine Engineering    70000  
## 5      Chemical Engineering    65000  
## 6      Nuclear Engineering    65000  
## 7      Actuarial Science    62000  
## 8 Astronomy And Astrophysics    62000  
## 9      Mechanical Engineering    60000  
## 10 Electrical Engineering    60000  
## # ... with 163 more rows
```

Summarize a data frame using `summarize()`

The average number of female grads and the total number of majors in the data set.

```
college_recent_grads %>%  
  select(major, men, women) %>%  
  summarise(femgrad_mean = mean(women), N = n())
```

```
## # A tibble: 1 x 2  
##   femgrad_mean      N  
##   <dbl> <int>  
## 1    22530.36    173
```

Summarize groups in a data frame using `summarize()` and `group_by()`

The median salary in majors with 45%-55% female students.

```
my_college_dat %>%  
  group_by(sex_equal) %>%  
  summarise(median(salary_median))
```

```
## # A tibble: 2 x 2  
##   sex_equal `median(salary_median)`  
##   <chr>          <dbl>  
## 1      No      37400  
## 2     Yes      36000
```

Combining Multiple Tables

Sentiment of Trump's Tweets

- Donald Trump likes to tweet a lot.
- Some tweets have an angry sentiment or contain insults, and some are not.
- Trump supposedly used to send tweets from a [Samsung Galaxy](#) when he is [insulting people, places, and things](#), from other devices such as an iPhone when he is not.
- Trump's last tweet from Android were March 25, 2017



Jonathan Ernst /

Is Trump Still Tweeting From His Unsecured Android Phone?

If so, he's extremely vulnerable to being hacked.

OFFER ENDS SOON

Get an ADT Interactive Video Package Installed for

NOW ONLY
\$24⁹⁹*

Trump's Tweets

A data frame with Trump's Tweets.


```
trumptweets <- read_csv("trumptweets.csv") #import from csv file
glimpse(trumptweets)
```

```
## Observations: 53,333
## Variables: 4
## $ source      <chr> "Android", "Android", "Android", "Android", "Androi...
## $ created_at  <dtm> 2013-02-06 01:53:40, 2013-02-06 01:53:40, 2013-02-...
## $ id_str      <dbl> 2.989727e+17, 2.989727e+17, 2.989727e+17, 2.989727e...
## $ word        <chr> "@sherrieshepherd", "nice", "comments", "view", "te..."
```


Trump's tweets

```
trumptweets %>%  
  count(word) %>%  
  mutate(word = reorder(word,n)) %>%  
  top_n(20) %>%  
  ggplot(aes(word, n)) + geom_col() + coord_flip() +  
  labs(x = "Word", y = "Number of times word occurs in a Tweet")
```

Word



Sentiment Lexicon

- Several lexicons (dictionaries) have been developed that categorize words according to sentiment (feeling or emotion).
- The `tidytext` library contains several lexicons.

```
library(tidytext)
sentiments
```

```
## # A tibble: 27,314 x 4
##       word sentiment lexicon score
##       <chr>      <chr>   <chr> <int>
## 1    abacus      trust    nrc    NA
## 2   abandon      fear    nrc    NA
## 3   abandon negative    nrc    NA
## 4   abandon      sadness  nrc    NA
## 5  abandoned      anger    nrc    NA
## 6  abandoned      fear    nrc    NA
## 7  abandoned negative    nrc    NA
## 8  abandoned      sadness  nrc    NA
## 9 abandonment      anger    nrc    NA
## 10 abandonment      fear    nrc    NA
## # ... with 27,304 more rows
```

NRC Lexicon

- The nrc lexicon categorizes words in a binary fashion (“yes”/“no”) into categories of positive, negative, anger, anticipation, disgust, fear, joy, sadness, surprise, and trust.
- The `getsentiments ()` function provides a way to get specific sentiment lexicons without the columns that are not used in that lexicon.

NRC Lexicon

```
get_sentiments("nrc")
```

```
## # A tibble: 13,901 x 2
##       word sentiment
##   <chr>      <chr>
## 1   abacus    trust
## 2  abandon    fear
## 3  abandon negative
## 4  abandon sadness
## 5 abandoned anger
## 6 abandoned fear
## 7 abandoned negative
## 8 abandoned sadness
## 9 abandonment anger
## 10 abandonment fear
## # ... with 13,891 more rows
```

Sentiment of Words used in Tweets

- To examine the sentiment of the words Trump used in tweets we need to join the data frame containing the NRC lexicon and the data frame of Trump's words used in tweets.
- `inner_join(x, y)`: return all rows from x where there are matching values in y, and all columns from x and y. If there are multiple matches between x and y, all combination of the matches are returned.

```
trumptweets %>% inner_join(get_sentiments("nrc"))
```

```
## # A tibble: 33,043 x 5
##   source      created_at    id_str      word sentiment
##   <chr>      <dtm>      <dbl>      <chr>      <chr>
## 1 Android 2013-02-06 01:53:40 2.989727e+17 terrific    sadness
## 2 Android 2013-02-18 23:36:36 3.036492e+17 sky         positive
## 3 Android 2013-02-18 23:36:36 3.036492e+17 rocket      anger
## 4 Android 2013-02-18 23:36:36 3.036492e+17 payback     anger
## 5 Android 2013-02-18 23:36:36 3.036492e+17 payback     negative
## 6 Android 2013-02-19 00:25:48 3.036616e+17 surprised   surprise
## 7 Android 2013-02-19 12:36:19 3.038455e+17 buss        joy
## 8 Android 2013-02-19 12:36:19 3.038455e+17 buss        positive
```

```
## 9 Android 2013-02-19 12:36:19 3.038455e+17 friend joy
## 10 Android 2013-02-19 12:36:19 3.038455e+17 friend positive
## # ... with 33,033 more rows
```

Sentiment of Words used in Tweets

```
## Joining, by = "word"
```


```
## # A tibble: 20 x 4
## # Groups:   sentiment [10]
##      sentiment source      n    pct
##      <chr>    <chr> <int> <dbl>
## 1      disgust Android  1537 80.68
## 2     negative Android  4040 78.68
## 3      sadness Android  2117 78.32
## 4        anger Android  2228 78.31
## 5         fear Android  2057 77.80
## 6     surprise Android  1297 72.70
## 7          joy Android  1777 71.65
## 8 anticipation Android  2240 71.25
## 9      positive Android  4328 69.72
## 10         trust Android  2924 69.70
## 11         trust  iPhone  1271 30.30
## 12      positive  iPhone  1880 30.28
## 13 anticipation  iPhone   904 28.75
## 14          joy  iPhone   703 28.35
## 15     surprise  iPhone   487 27.30
## 16         fear  iPhone   587 22.20
## 17         anger  iPhone   617 21.69
## 18      sadness  iPhone   586 21.68
## 19     negative  iPhone  1095 21.32
## 20      disgust  iPhone   368 19.32
```


Sentiment of Words used in Tweets

Now let's plot our results

```
trumptweets %>%  
  inner_join(get_sentiments("nrc")) %>%  
  group_by(sentiment, source) %>%  
  count(sentiment) %>%  
  ungroup() %>%  
  mutate(sentiment = reorder(sentiment, n)) %>%  
  group_by(sentiment, source) %>%  
  ggplot(aes(sentiment, n)) + geom_col(aes(fill = source), position = "dodge") + coord_flip() +  
  labs(y = "Number of words in Tweets", x = "NRC Sentiment")
```

NRC Sentiment



Join two tables together

- In the dplyr library there are several other ways to join tables:
`left_join()`, `right_join()`, `full_join()`, `semi_join()`,
`anti_join()` .
- See dplyr [documentation](#).

Transforming data

Statistical Transformations

- In statistical analysis it's often necessary to transform data.
- Transforming data takes each value of a variable x_i and transforms it into $f(x_i)$:


$$x_i \mapsto f(x_i).$$

- Common transformations include: $f(x) = \ln(x)$, and $f(x) = x^p$, $p \in \mathbb{R}$. For example, if $p = 1/2$ then f is the square-root transformation.

Logarithmic transformation

- Logarithmic transformation refers to the natural logarithm:


$$y = \log_e(x) \iff \exp(y) = e^y = x$$



Transforming Variables in R

The relationship between Salary (median) and percentage of male graduates.

```
college_recent_grads %>%  
  ggplot(aes(x = men / total, y = median)) + geom_point()
```

```
college_recent_grads %>%  
  mutate(log_men = log(men / total), log_salary = log(median)) %>%  
  ggplot(aes(x = log_men, y = log_salary)) + geom_point()
```

